



Last updated on **Nov 5, 2025**

Compute

Databricks compute refers to the selection of computing resources available on Databricks to run your data engineering, data science, and analytics workloads. Choose from serverless compute for on-demand scaling, classic compute for customizable resources, or SQL warehouses for optimized analytics.

You can view and manage compute resources in the **Compute** section of your workspace:

Serverless compute

On-demand, automatically managed compute that scales based on your workload requirements.

Topic	Description
Serverless compute for notebooks	Interactive Python and SQL execution in notebooks with automatic scaling and no infrastructure management.
Serverless compute for jobs	Run Lakeflow Jobs without configuring or deploying infrastructure. Automatically provisions and scales compute resources.
Serverless pipelines	Run Lakeflow Spark Declarative Pipelines without configuring or deploying infrastructure. Automatically provisions and scales compute resources.

Topic	Description
Serverless compute limitations	Understanding limitations and requirements for serverless workloads and supported configurations.

Classic compute

Provisioned compute resources that you create, configure, and manage for your workloads.

Topic	Description
Classic compute overview	Overview of who can access and create classic compute resources.
Configure compute	Create and configure compute for interactive data analysis in notebooks or automated workflows with Lakeflow Jobs.
Standard compute	Multi-user compute with shared resources for cost-effective collaboration. Lakeguard provides secure user isolation.
Dedicated compute	Compute resource assigned to a single user or group.
Instance pools	Pre-configured instances that reduce compute startup time and provide cost savings for frequent workloads.

SQL warehouses

Optimized compute resources for specific use cases and advanced functionality. SQL warehouses can be configured as serverless or classic.

Topic	Description
SQL warehouses	Optimized compute for SQL queries, analytics, and business intelligence workloads with serverless or classic options.
SQL warehouse types	Understanding the differences between serverless and classic SQL warehouse options to choose the right type for your workloads.

Additional topics

Topic	Description
What is Photon?	High-performance query engine that accelerates SQL workloads and provides faster data processing.
What is Lakeguard?	Security framework that provides data governance and access control for compute resources.

For information about working with compute using the command line or APIs, see [What is the Databricks CLI?](#) and the [Databricks REST API reference](#).

Reserved ports

Certain ports are reserved on the driver node for internal Databricks services. To avoid conflicts, do not bind services to the following ports:

- 1023
- 6059
- 6060
- 6061
- 6062: Occupied by ipywidgets by default. You can change the port if needed. See [ipywidgets](#).
- 7071
- 7077
- 10000
- 15001
- 15002
- 36423
- 38841
- 39909
- 40000
- 40001
- 41063